

## Seminar Series Neural Networks for Finance

### Lecture 2

#### Aim

Neural networks in depth: capacity and optimization

24-02-2022

Balint Negyesi ([b.negyesi@tudelft.nl](mailto:b.negyesi@tudelft.nl))

*Deep versus Shallow slides written by Nikolaj Mücke  
([nikolaj.mucke@cwi.nl](mailto:nikolaj.mucke@cwi.nl))*

# Agenda

- 1 Neural networks: ANNs, DNNs, activations
- 2 Universal Approximation Theorem
- 3 Stochastic Gradient Descent
- 4 Automatic differentiation: descending into a computer
- 5 Gradient based optimization problems
- 6 Deep versus Shallow
  - Shallow Networks

# (Fully-Connected) Feedforward Artificial Neural Networks

- A **fully-connected, feedforward artificial neural network** is simply a **hierarchical composition** of the transformations above, mapping an input  $x \in \mathbb{R}^d$  to the output  $\Phi \in \mathbb{R}^q$  whose  $i$ 'th coordinate reads as follows

$$\Phi_i(x|\theta) := \varphi_i^{(2)} \left( \sum_{j=1}^{p_1} W_{ij}^{(2)} \varphi_j^{(1)} \left( \sum_{k=1}^d W_{jk}^{(1)} x_k + b_j^{(1)} \right) + b_i^{(2)} \right) \quad (1)$$

where  $\theta := (\theta^{(1)}, \theta^{(2)}) := (W^{(1)}, b^{(1)}, W^{(2)}, b^{(2)}) \in \mathbb{R}^{p_1 \times (d+1) + q \times (p_1+1)}$

- The total number of parameters in the statistical model is

$$p = p_1 \times (d + 1) + q \times (p_1 + 1) \quad (2)$$

## Disclaimer


From this lecture on – as is usual in the field of statistics (and also machine learning) – we drop the bold face notation for vectors and matrices (and tensors). Tensor valued entries should be interpreted by context. For example:  $x \leftarrow \mathbf{x} \in \mathbb{R}^d$ ,  $\Theta \leftarrow \mathbf{\Theta} \in \mathbb{R}^p$ ,  $W^{(l)} \leftarrow \mathbf{W}^{(l)} \in \mathbb{R}^{p_l \times p_{l-1}}, \dots$

# (Fully-Connected) Feedforward Artificial Neural Networks


- Denoting the *element-wise* non-linearities by  $\varphi^{(n)}(z^{(n)}) := (\varphi_1^{(n)}(z_1^{(n)}), \dots, \varphi_{p_n}^{(n)}(z_{p_n}^{(n)})) \in \mathbb{R}^{p_n}$ , in vector notation this can be written as follows

$$\Phi(x|\Theta) := \varphi^{(2)} \circ z^{(2)}(\cdot|\theta^{(2)}) \circ \varphi^{(1)} \circ z^{(1)}(x|\theta^{(1)}) \quad (3)$$

- FCFF ANNs can be thought of as **directed, acyclic graphs**




## Example: XOR revisited



- Unlike with SLP, adding a *hidden layer*, and parametrizing with ANNs solves the problem
- Let  $h_1(z_1)$  and  $h_2(z_2)$  be perceptrons with  $b_1 = b_2 = 0$
- Fix  $w_1 = -1, w_2 = 1, w_3 = 1, w_4 = -1$
- The desired output is obtained by setting  $w_5 = w_6 = 1$  and  $b_3 = 1$

## Example: XOR revisited



- $\hat{\Theta} = (w_1 = -1, w_2 = 1, w_3 = 1, w_4 = -1, b_1 = 0, b_2 = 0, w_5 = 1, w_6 = 1, b_3 = 1)$  gives and the XOR problem is solved

$x_1$	$x_2$	$z_1$	$z_2$	$y$
-1	-1	1	1	1
-1	1	1	-1	-1
1	-1	-1	1	-1
1	1	-1	-1	1

Activations: the key feature of non-linearity



# Activation Functions


- The key component of *non-linearity* is given by the **activation functions**  $\varphi_j^{(n)} : \mathbb{R} \rightarrow \mathbb{R}, j = 1, \dots, p_n$
- For SLP we had  $\varphi_1^{(1)}(x) = \mathbb{1}_{x-b>0}(x)$
- Most often, application functions within the same layer are chosen the same  $\varphi_1^{(n)}(x) = \dots = \varphi_{p_n}^{(n)}(x)$
- In regression applications the output activations are usually chosen to be the identity function  $\varphi^{(2)}(z^{(2)}) = z^{(2)}$
- The parametrization inherits **continuity** and **differentiability** properties from the chosen activations
- The non-linear activation functions are *hyperparameters* of the statistical model, appropriate choices vary depending on the problem

# Activation Functions

name	$\varphi(x)$	$\varphi'(x)$	range	continuity
heaviside	$\begin{cases} 0, & x \leq 0 \\ 1, & x > 0 \end{cases}$	$\begin{cases} 0, & x \neq 0 \\ \text{undefined}, & x = 0 \end{cases}$	$\{0, 1\}$	$C^{-1}$
ReLU	$\begin{cases} 0, & x \leq 0 \\ x, & x > 0 \end{cases}$	$\begin{cases} 0, & x < 0 \\ x, & x > 0 \\ \text{undefined}, & x = 0 \end{cases}$	$[0, \infty)$	$C^0$
Gaussian	$e^{-x^2}$	$-2xe^{-x^2}$	$(0, 1]$	$C^\infty$
tanh	$\frac{e^x - e^{-x}}{e^x + e^{-x}}$	$1 - \tanh^2(x)$	$(-1, 1)$	$C^\infty$
logistic	$\frac{1}{1 + e^{-x}}$	$\frac{e^{-x}}{1 + e^{-x}}$	$(0, 1)$	$C^\infty$
softplus	$\log(1 + e^x)$	$\frac{1}{1 + e^{-x}}$	$(0, \infty)$	$C^\infty$
...	...	...	...	...

See wiki

# Activation Functions



# Layer dependent activations

- For classification problems it is chosen such that  $\sum_{k=1}^q \varphi_k^{(2)}(z_k^{(2)}) = 1$
- In order to ensure that output *probabilities* indeed sum to one, activations are chosen such that they are not a collection of scalar-valued mappings of neurons but vector-valued functions depending on the whole layer  $\varphi^{(n)} : \mathbb{R}^{p_n} \rightarrow \mathbb{R}^{p_n}$

name	$\varphi_i(x)$	$\partial_j \varphi_i(x)$	range	cont.
softmax	$\frac{e^{x_i}}{\sum_{j=1}^{p_n} e^{x_j}}$	$\varphi_i(x)(\delta_{ij} - \varphi_j(x))$	$(0, 1)$	$C^\infty$
maxout	$\max_{i=1, \dots, p_n} x_i$	$\begin{cases} 1, & j = \arg \max_i x_i \\ 0, & j \neq \arg \max_i x_i \end{cases}$	$(-\infty, \infty)$	$C^0$
...	...	...	...	...

## Deep Neural Networks

# Fully-Connected Feedforward Deep Neural Networks

- So far: neural network = input + hidden + output layer
- These are called **shallow** feedforward ANNs

$$\Phi(x|\Theta) := \varphi^{(2)} \circ z^{(2)}(\cdot|\theta^{(2)}) \circ \varphi^{(1)} \circ z^{(1)}(x|\theta^{(1)}) \quad (4)$$

- One can extend the model to allow for  $L$ -many hidden layers in the composition, with corresponding widths  $p_l, l = 1 \dots, L$
- The resulting, **deep neural network** mapping reads as follows

$$\Psi(x|\Theta) := \varphi^{(L+1)} \circ z^{(L+1)}(\cdot|\theta^{(L+1)}) \circ \dots \circ \varphi^{(1)} \circ z^{(1)}(x|\theta^{(1)}), \quad (5)$$

where  $\Theta := (\theta^{(1)}, \dots, \theta^{(L+1)}) \in \mathbb{R}^p$

- The total number of parameters in the statistical model rapidly increases

$$p = d \times (p_1 + 1) + \sum_{l=1}^{L-1} p_l \times (p_{l+1} + 1) + p_L \times (q + 1) \quad (6)$$

# Fully-Connected Feedforward Deep Neural Networks

$$\Psi(x|\Theta) := \varphi^{(L+1)} \circ z^{(L+1)}(\cdot|\theta^{(L+1)}) \circ \dots \circ \varphi^{(1)} \circ z^{(1)}(x|\theta^{(1)}), \quad (7)$$




Figure: Illustration DNN, source: NN-SVG

$$p = d \times (p_1 + 1) + \sum_{l=1}^{L-1} p_l \times (p_{l+1} + 1) + p_L \times (q + 1)$$

## Universal Approximation Theorem



# Most commonly used carpet

## Topic

- We have seen: (deep) neural networks provide *one way* to parametrize non-linear mappings
- We have not seen: **is that a "good" way? What class of functions can actually be approximated by them?**

# A classical analogy

## Theorem (Weierstrass approximation theorem)

Let  $f \in C([a, b]; \mathbb{R})$ . For any  $\epsilon > 0$  there exists a polynomial  $p : x \mapsto a_0 + a_1x + a_2x^2 + \dots$  such that  $\sup_{x \in [a, b]} |f(x) - p(x)| < \epsilon$ .

- **Weierstrass approximation theorem:** any continuous function can be uniformly approximated by a polynomial functions over closed intervals, with arbitrary accuracy
- Terminology the set of polynomials is **dense** in  $C([a, b]; \mathbb{R})$
- Conclusion: polynomials may be a good parametrization for continuous functions
- Notice: **not implementable**
  - 1 No bound on the order of the polynomial
  - 2 No way to find the "right" polynomial within the class

# Universal Approximation Theorem

- In order to show that neural networks are "good" family of statistical models, we need a similar result → **Universal Approximation Theorem (UAT)**
- First: Weierstrass for shallow neural networks
- Definition: an activation  $\varphi : \mathbb{R} \rightarrow \mathbb{R}$  is called **sigmoidal**, if it is continuous and  $\lim_{x \rightarrow -\infty} \varphi(x) = 0, \lim_{x \rightarrow \infty} \varphi(x) = 1$

## Theorem (UAT – continuous functions, Cybenko, 1989)

Let  $\mathcal{G}_\varphi := \{\Phi(x|\Theta) : \Theta \in \mathbb{R}^p, p \in \mathbb{N}, \varphi \text{ fixed, sigmoidal}\}$  be the family of shallow neural networks of the form (3) with  $d = q = 1$ . Then, for any  $f \in C([0, 1]; \mathbb{R})$  and  $\epsilon > 0$  there exists a shallow neural network  $\Phi(\cdot|\Theta^*) \in \mathcal{G}_\varphi$  such that

$$\sup_{x \in [0,1]} |\Phi(x|\Theta^*) - f(x)| < \epsilon.$$

# Restrictions and extensions

- Similar restrictions as in Weierstrass: *arbitrary width*, no insight on  $\Theta^*$
- Natural extension from shallow to deep neural networks of the form (7)
- Several results exist, very actively researched area since the early 90s
  - 1 sigmoidal  $\rightarrow$  non polynomial
  - 2  $[a, b] \rightarrow \mathbb{R} \rightarrow$  compact subsets in general vector spaces to other vector spaces
  - 3 dense in  $C([a, b]; \mathbb{R}) \rightarrow$  dense in the space of Lebesgue integrable functions ( $L^p$  spaces)
  - 4 approximating functions with arbitrary accuracy  $\rightarrow$  approximating functions and **their derivatives** with arbitrary accuracy (**Sobolev spaces**) – see Hornik et al., 1990
  - 5 arbitrary width/depth  $\rightarrow$  fixed width/depth with dominated error term (**curse of dimensionality**)
  - 6 ...

## Stochastic Gradient Descent

## Question

- UAT  $\rightarrow$  neural networks may be an appropriate statistical model for a wide range of phenomena
- UAT: in the space of neural networks there lives a neural network  $\Phi(x|\Theta^*)$  which approximates a given function  $f(x)$  with arbitrary accuracy
- We do not know: **how do we find  $\Theta^*$ , or at least a good approximation of it?**

# Iterative schemes

- Recall: thanks to the **convexity** of  $\beta \mapsto \|y - X\beta\|^2$ , OLS has a closed form expression mapping measurements  $X, y$  to the **unique, global, minimizer** of MSE

$$\beta^* = (X^T X)^{-1} X^T y \quad (8)$$

- In neural network regression the loss is no longer convex  $L(\Theta|D = (x, y)) = \|y - \Phi(x|\Theta)\|^2$  due to the non-linearity of  $\Phi(x|\Theta)$  in  $\Theta$
- No closed-form expression for a minimizer  $\implies$  **iterative schemes** needed
- No unique solution** to  $\nabla_{\Theta} L(\Theta|D) = 0$
- A solution to  $\nabla_{\Theta} L(\Theta|D) = 0$  may not be an extremum at all (**saddle points**)
- local minima** may not be global  $\implies$  solving  $\nabla_{\Theta} L(\Theta|D) = 0$  does not guarantee global minimum

# Iterative schemes – general formulation

- An iterative scheme is a *sequence of approximations*  $\Theta^n$  such that the resulting approximations  $\Phi^n := \Phi(\cdot | \Theta^n)$  converge to a desired solution  $H$  in an appropriate sense (e.g. MSE)
- An iterative scheme needs an appropriate *update/iterative rule* which determines  $\Theta^{n+1}$  given  $\Theta^n$
- How do we define a sensible update rule?  $\longrightarrow$  **(Stochastic) Gradient Descent ((S)GD)**



# Gradient Descent – intuition

- Assume we want to minimize some  $f : \mathbb{R}^p \rightarrow \mathbb{R}_+$  non-linear real-valued function with some iterative scheme. Our current guess for the minimum is  $x^k$ . **What should our next guess be?**  
→ **steepest gradient**
- It is always nice to gather an intuition in one dimension




Figure: Goodfellow et al.: Deep Learning (2016), MIT Press [Figure 4.1, pg. 81]

# Gradient Descent – general

- Back to the general setting: given a point  $x^k \in \mathbb{R}^p$  we want to find the **direction**  $e \in \mathbb{R}^p$ ,  $\|e\| = 1$  in which the value of  $f$  decreases the fastest
- Formally we are concerned with the following minimization problem

$$\begin{aligned}\arg \min_{e \in \mathbb{R}^p, \|e\|=1} \left\langle e \mid \nabla_x f(x^k) \right\rangle &= \arg \min_{e \in \mathbb{R}^p, \|e\|=1} \|e\| \left\| \nabla_x f(x^k) \right\| \cos(\vartheta_{e,x^k}) \\ &= \arg \min_{e \in \mathbb{R}^p, \|e\|=1} \left\| \nabla_x f(x^k) \right\| \cos(\vartheta_{e,x^k}) \\ &= \arg \min_{e \in \mathbb{R}^p, \|e\|=1} \cos(\vartheta_{e,x^k}),\end{aligned}$$

where  $\vartheta_{e,x^k}$  is the angle between  $e$  and  $\nabla_x f(x^k)$

- This is minimized when  $\vartheta_{e,x^k} = \pi$  leading to  
 $e = -\nabla_x f(x^k) / \left\| \nabla_x f(x^k) \right\| \implies$  **steepest descent** or  
**gradient descent**

## Gradient Descent (GD)

The iterative scheme of a gradient descent optimization then reads as

$$x^{k+1} = x_k - \eta^k \nabla_x f(x^k), \quad (9)$$

where the parameter  $\eta^k$  is called a **learning rate** and determines the size of each iteration step

Several strategies for the choice of  $k \rightarrow \eta^k$ : constant, adaptive (decay, validation loss linked, etc.), line search, ...

# Back to machine learning (statistics)

Our goal is to minimize a performance measure (loss)

$$L(\Theta) := \mathbb{E}_{(x,y) \sim \mathbb{P}_{\text{data}}} [d(\Phi(x|\Theta), y)] \quad (10)$$

**indirectly** by minimizing an empirical loss defined by

$$L(\Theta|D) := \frac{1}{M} \sum_{m=1}^M d(\Phi(x_m|\Theta), y_m) \quad (11)$$

given a finite number of  $M$  measurements drawn from the same data generating distribution

**capacity, generalization, overfitting, ...**

# From Gradient Descent to Stochastic Gradient Descent

- The GD iteration in (9) on (11) would read as

$$\Theta^{k+1} = \Theta^k - \eta^k \frac{1}{M} \sum_{m=1}^M \nabla_{\Theta} d(\Phi(x_m | \Theta^k), y_m) \quad (12)$$

- This iteration scheme is sometimes called **batch gradient descent** (*personal opinion: wrongly*)
- Drawbacks: computationally intensive (gradient for every sample), large memory requirements, cannot update model on the fly, ...
- Idea: random partition of the dataset and obtain gradient steps on these random subsets → **Stochastic Gradient Descent (SGD)**

# Stochastic Gradient Descent

The steepest descent is decomposed into steps. Choose  $B < M$ .

Choose a random permutation of the dataset

$\pi : \{1, \dots, M\} \rightarrow \{1, \dots, M\}$ .

- 1 Random permutation of the dataset into  $\pi^k$
- 2 Loop over each resulting subset  $i = 1, \dots, \lceil M/B \rceil$  of size  $B$ :  
 $\{x_{\pi_m}, y_{\pi_m}\}_{m=(i-1)B}^{\min(iB, M)}$
- 3 Update the parameters over each subset  $i = 1, \dots, \lceil M/B \rceil$

$$\Theta^{k,i+1} = \Theta^{k,i} - \eta^k \frac{1}{B} \sum_{m=(i-1)B}^{\min(iB, M)} d(\Phi(x_m | \Theta^{k,i}), y_m) \quad (13)$$

- 4 The  $k + 1$ 'th iteration step receives  
 $\Theta^{k+1} = \Theta^{k+1,1} := \Theta^{k, \lceil M/B \rceil}$

Terminology:  $k$ : **epoch**,  $i$ : **batch**,  $B$ : **batch size**

# Batch size

- $B = 1$  is often called **online learning**
- Larger  $B$ s provide more accurate estimates of the gradient  $\nabla_{\Theta} L(\Theta|D)$  but with diminishing returns
- Small batches often have a *regularization effect* and thus generalize better
- Small batches require more iterations per epoch and are thus slower
- Eventually: **trial and error...**

# Initialization and convergence criteria


Any implementable iterative optimization scheme requires two more ingredients

- An **initialization**  $\Theta^0$  determining the sequence  $\Theta^k \leftarrow \Theta^{k-1}, k \geq 1$ 
  - biases are usually initialized with an arbitrarily chosen constant value – most often 0
  - weights are initialized randomly, several options exist:  
 $W_{ij}^{(l)} \sim U\left(-\sqrt{\frac{6}{p_{l-1}+p_l}}, \sqrt{\frac{6}{p_{l-1}+p_l}}\right), W_{ij}^{(l)} \sim \mathcal{N}\left(0, \sqrt{\frac{2}{p_{l-1}+p_l}}\right)$
  - usually designed in a way to fix the variances of each layer's output activations, but also the variances of the gradients
  - **crucial step** for any successful iteration (no local optima, ...), actively researched topic
- A **convergence criterion** which stops the iteration at a given step  $K$  concluding in approximations  $\hat{\Theta} := \Theta^K$  – based on validation loss, maximum iteration number, ...



# Learning rate

- The size of an iteration step in (9)–(13) in the *steepest direction* depends on the **learning rate**
- Closer to the minimum, gradients tend to decrease; smaller steps can thus lead to a better approximation of the minimum
- Several strategies  $k \rightarrow \eta^k$ : constant, adaptive (decay, validation loss linked, etc.), line search, ...
- Eventually: **trial and error**
- **Momentum** based estimates: momentum SGD, Nesterov, ...
- Alternative **adaptive** optimization methods: Adam, AdaGrad, RMSProp, ...



Automatic differentiation: descending into a computer


- We have seen: gradient based iterative optimization methods to minimize empirical loss
- However, they by definition rely on being able to calculate the **gradient** of the loss with respect to the parameters in the model  $\nabla_{\Theta} L(\Theta|D)$
- If  $d(f, g)$  in 11 is the mean-squared error we, e.g., have  $\nabla_{\Theta} L(\Theta|D) = 2(\nabla_{\Theta} \Phi(x|\Theta))^T (\Phi(x|\Theta) - y)$
- $\Theta \in \mathbb{R}^p$  is a very high-dimensional vector (huge parameter spaces) – (6)

## Question

How can such a large dimensional gradient can be computed in a **precise, efficient, and robust** way?

Finite differences: high errors, slow, bad scaling; symbolic expressions: memory consumption, very long and redundant expressions  $\longrightarrow$  **automatic differentiation**

# Recall: neural networks' hierarchical structure



- Hierarchical sequence of compositions of the form (7)
- Complete parameter space  
 $\Theta := (W^{(1)}, b^{(1)}, \dots, W^{(L+1)}, b^{(L+1)}) \in \mathbb{R}^p$  with  
 $W^{(l)} \in \mathbb{R}^{p_l \times p_{l-1}}$ ,  $b^{(l)} \in \mathbb{R}^{p_l}$  and  $p_0 = d$
- Activations by non-linear mappings  $\varphi^{(l)}$  and their affine combinations

$$z^{(l+1)} = W^{(l+1)}z^{(l)} + b^{(l+1)}, \quad a^{(l+1)} = \varphi^{(l+1)}(z^{(l+1)}),$$

$$l = 0, \dots, L$$

# Computational graphs

- This structure

$$x \xrightarrow{W^{(1)}, b^{(1)}} z^{(1)} \xrightarrow{\varphi^{(1)}} a^{(1)} \mapsto \dots \xrightarrow{\varphi^{(L+1)}} a^{(L+1)}$$

can be represented as **directed, acyclic graph**  $\longrightarrow$   
**computational graph**

- Each node in the graph corresponds to an operation and *data flows through the vertices*
- An operation is characterized by three characteristics: a compute function which determines the node's output; set of input nodes; and set of output nodes
- Computational graphs: complex calculations decomposed into a sequence of elementary operations; far more general than machine learning

# Forward propagation




Figure: Forward Propagation

- The input nodes of  $z^{(l)}$  are  $a^{(l-1)}, W^{(l)}, b^{(l)}$

$$z_i^{(l)} = \sum_{j=1}^{p_{l-1}} W_{ij}^{(l)} a_j^{(l-1)} + b_i^{(l)} \quad (14)$$

- With respect to these inputs are easily obtained

$$\frac{\partial z_i^{(l)}}{\partial a_j^{(l-1)}} = W_{ij}^{(l)}, \quad \frac{\partial z_i^{(l)}}{\partial b_j^{(l)}} = \delta_{ij}, \quad \frac{\partial z_i^{(l)}}{\partial W_{kj}^{(l)}} = \delta_{ki} a_j^{(l-1)} \quad (15)$$

- The input nodes of  $a^{(l)}$  are in  $z^{(l)}$ :  $a_i^{(l)} = \varphi_i^{(l)}(z^{(l)})$

$$\frac{\partial a_i^{(l)}}{\partial z_j^{(l)}} = (\nabla_{z^{(l)}} \varphi^{(l)}(z^{(l)}))_{ij} \quad (16)$$

# Derivation

With Jacobian matrices

$$\begin{array}{l} \text{Forward} \\ \text{Backward} \end{array} \left\{ \begin{array}{ll} a^{(0)} & = x, \\ z^{(1)} & = W^{(1)} \cdot a^{(0)} + \mathbf{b}^{(1)}, \\ a^{(1)} & = \varphi^{(1)}(z^{(1)}), \\ \dots & \\ z^{(L+1)} & = W^{(L+1)} \cdot a^{(L)} + \mathbf{b}^{(L+1)}, \\ a^{(L+1)} & = \varphi^{(L+1)}(z^{(L+1)}) \end{array} \right. \quad (17)$$
$$\left\{ \begin{array}{ll} \nabla_x a^{(0)} & = I_d, \\ \nabla_{a^{(0)}} z^{(1)} & = W^{(1)}, \\ \nabla_{z^{(1)}} a^{(1)} & = \nabla_{z^{(1)}} \varphi^{(1)}(z^{(1)}), \\ \dots & \\ \nabla_{a^{(L)}} z^{(L+1)} & = W^{(L+1)}, \\ \nabla_{z^{(L+1)}} a^{(L+1)} & = \nabla_{z^{(L+1)}} \varphi^{(L+1)}(z^{(L+1)}) \end{array} \right.$$



# Back-propagation

- Suppose there is one node in the last layer, matrix  $W^{(L+1)} = [w_{1j}^{(L+1)}, j = 1, \dots, p_L]$  is a vector then
- We compute the sensitivity of the mean-squared distance function  $d$  with respect to weight  $w_{1j}^{(L+1)}$  by the **chain rule**

$$\begin{aligned}\frac{\partial d}{\partial W_{1j}^{(L+1)}} &= \frac{\partial d}{\partial a_1^{(L+1)}} \frac{\partial a_1^{(L+1)}}{\partial W_{1j}^{(L+1)}} \\ &= \frac{\partial d}{\partial a_1^{(L+1)}} \frac{\partial a_1^{(L+1)}}{\partial z_1^{(L+1)}} \frac{\partial z_1^{(L+1)}}{\partial W_{1j}^{(L+1)}} \\ &= 2(a_1^{(L+1)} - y) \varphi^{(L+1)'}(z_1^{(L+1)}) a_j^{(L)},\end{aligned}$$

from Eq. (15)

- Similarly, the **chain rule** gives an expression for the derivative of  $L$  wrt the inputs  $x$

$$\begin{aligned}(\nabla_x d)^T &= (\nabla_{a^{(L+1)}} d)^T \nabla_{z^{(L+1)}} a^{(L+1)} \nabla_{a^{(L)}} z^{(L+1)} \dots \quad (18) \\ &\quad \dots \nabla_{z^{(1)}} a^{(1)} \nabla_{a^{(0)}} z^{(1)} \nabla_x a^{(0)} \\ &= (\nabla_{a^{(L+1)}} d)^T \nabla_{z^{(L+1)}} \varphi^{(L+1)}(z^{(L+1)}) W^{(L+1)} \dots \\ &\quad \dots \nabla_{z^{(1)}} \varphi^{(1)}(z^{(1)}) W^{(1)} I\end{aligned}$$

where each element is a Jacobian matrix

- **Three types of terms:**, the gradient of the distance with respect to the outputs  $\nabla_{a^{(L+1)}} d$ ; Jacobian matrix of each layer's output activations  $\nabla_{z^{(l)}} \varphi^{(l)}(z^{(l)})$ ; and the derivative of the affine transformation function with respect to input activations  $\nabla_{a^{(l-1)}} z^{(l)}$

# Backpropagation

To derive the closed-form expressions of backward propagation, we define the auxiliary error vector

$$[\delta^{(l)}]^T := [\nabla_{a^{(L+1)}} d]^T \nabla_{z^{(L+1)}} \varphi^{(L+1)}(z^{(L+1)}) W^{(L+1)} \dots \quad (19) \\ \dots W^{(l+1)} \nabla_{z^{(l)}} \varphi^{(l)}(z^{(l)}),$$

which is a vector of length  $\mathbb{R}^{p_l}$

## Backpropagation

The backpropagation scheme, implementing the derivative of the loss function is then defined by the following recursive expressions

$$\begin{aligned} \nabla_{b^{(l)}} L &= \delta^{(l)}, \\ \nabla_{W^{(l)}} L &= \delta^{(l)} (z^{(l-1)})^T, \\ [\delta^{(l)}]^T &= [\delta^{(l+1)}]^T (W^{(l+1)}) \nabla_{z^{(l)}} \varphi^{(l)}(z^{(l)}). \end{aligned} \quad (20)$$

# Backward propagation




Figure: Backpropagation

Gradient based optimization problems

# Challenges: local minima, saddle points, etc.

Together with backpropagation/automatic differentiation, SGD is a fully-implementable iterative optimization algorithm – their combination is called **training**

However, not without challenges

- local minima
- saddle points
- cliffs and plateaus: exploding and vanishing gradients
- ill-conditioned problems
- inaccurate gradient estimates – over batches
- regularization: **overfitting**
- . . .

# Local minima




Figure: Goodfellow et al.: Deep Learning (2016), MIT Press [Figure 4.3, pg. 83]

# Saddle points




Figure: Goodfellow et al.: Deep Learning (2016), MIT Press [Figure 4.5, pg. 88]



# Vanishing gradients, plateaus, cliffs, etc.

- The gradient in the current layer is equal to a product of terms of all the later layers
- Vanishing gradient problem when one term is close to zero; gradient explosion when the product approaches infinity  
(numerical instabilities, rounding errors)




Figure: Goodfellow et al.: Deep Learning (2016), MIT Press [Figure 8.3, pg. 285]

Remedies: **ReLU** – unbounded from above; renormalization; gradient clipping, etc.

Without the sake of completeness

- Input normalization
- **Batch normalization**: reparametrization, shifts and renormalizes activations within a layer across batches, *parametrizes* appropriate the appropriate means and variances
- **Dropout** is technique, which randomly turns off certain neurons (i.e., the gradient becomes zero for those neurons) to force the ANN to learn more features rather than saturate
- Constrained optimization: regularization with  $L^1 - L^2$  penalties, etc.
- ...

# Curvature, adaptive and second-order optimization

- Momentum: descent modeled as a particle's momentum with unit mass; previous updates' gradients decay with some hyperparameter  $\alpha$

$$\begin{aligned}v^{k+1} &= \alpha v^k - \eta^k \nabla_{\Theta} L(\Theta^k | D) \\ \Theta^{k+1} &= \Theta^k + v^{k+1}\end{aligned}\tag{21}$$

- Second-order optimization schemes: **Newton's method** to *approximate* the root of a multivariate function  $g : \mathbb{R}^n \rightarrow \mathbb{R}^n$  by

$$z^0 \in \mathbb{R}^n, \quad z^{k+1} = z^k - [\nabla_z g(z^k)]^{-1} g(z^k), \tag{22}$$

where  $\nabla_z g(z)$  is the Jacobian matrix of the vector-valued mapping  $g$  defined by  $[\nabla_z g(z)]_{ij} := \partial_{z_j} g_i(z)$

# Curvature, adaptive and second-order optimization




Figure: Goodfellow et al.: Deep Learning (2016), MIT Press [Figure 8.5, pg. 293]


Deep versus Shallow

# Deep versus Shallow

## Questions


- When does it make sense to go deep?
  - Why does it make sense to go deep?
- 
- Poggio, T., Mhaskar, H., Rosasco, L., Miranda, B., & Liao, Q. (2017). Why and when can deep – but not shallow – networks avoid the curse of dimensionality: A review. *arXiv:1611.00740 [cs]*
  - Mhaskar, H., Liao, Q., & Poggio, T. (2017). When and why are deep networks better than shallow ones? [Number: 1]. *Proceedings of the AAAI Conference on Artificial Intelligence*, 31(1)
  - Mhaskar, H., & Poggio, T. (2016). Deep vs. shallow networks : An approximation theory perspective. *arXiv:1608.03287 [cs, math]*

# Empirical difference



RMSE = 0.2449

**1 hidden layer** with 302 neurons  
Total number of neurons: 302



RMSE = 0.0014

**3 hidden layers** with 3 neurons  
Total number of neurons: 25

# Degree of Approximation

- **Complexity:** number of parameters  $\sim$  of neurons/units in a network
- Let  $\mathcal{V}_N := \{\Phi(\cdot|\Theta) : \Theta \in \mathbb{R}^N\} \subseteq \mathcal{G}$  be the set of networks with complexity  $N$ ,  $\mathcal{V}_N \subseteq \mathcal{V}_{N+1}$
- **Degree of Approximation** is defined by

$$\text{dist}(f, \mathcal{V}_N) := \inf_{\Phi \in \mathcal{V}_N} \|f - \Phi\|_X,$$

for  $f \in \mathcal{X}$



- $\text{dist}(f, \mathcal{V}_N) = \mathcal{O}(N^{-\gamma}) \implies$  a network with complexity  $N = \mathcal{O}(\epsilon^{-\frac{1}{\gamma}})$  is sufficient to guarantee accuracy at least  $\epsilon$



# Compositional Functions

We will consider **hierarchical compositions** of functions, such as

$$f(x_1, \dots, x_8) = h_3 \left( \begin{array}{l} h_{21}(h_{11}(x_1, x_2), h_{12}(x_3, x_4)), \\ h_{22}(h_{13}(x_5, x_6), h_{14}(x_7, x_8)) \end{array} \right)$$



# Compositional Functions' Class

- Let  $I^n = [-1, 1]^n$  and  $C(I^n)$  be the space of continuous functions on  $I^n$  with norm

$$\|f\| = \max_{x \in I^n} |f(x)|$$

- Then consider the subspace  $\mathcal{W}_r^n \subset C(I^n)$  consisting of  $r$  times continuously differentiable functions on  $I^n$ , where


$$\|f\| + \sum_{1 \leq |k|_1 \leq r} \|D^k f\| \leq 1,$$

where  $k$  is a multiindex and  $D^k$  denotes the partial derivative indicated by  $k$ , and  $|k|_1$  is the sum of the multiindex

# Compositional Functions - Function Class

- Let  $\mathcal{W}_r^{n,2}$  be the class of functions with a hierarchical structure with each constituent function,  $h_{ij} \in \mathcal{W}_r^2$ 
  - Binary tree functions
- Note  $\mathcal{W}_r^{n,2} \subset \mathcal{W}_r^n$
- Example of a  $\mathcal{W}_r^{4,2}$  function:

$$f(x_1, x_2, x_3, x_4) = \underbrace{h_2}_{\in \mathcal{W}_r^2} \left( \underbrace{h_{11}}_{\in \mathcal{W}_r^2}(x_1, x_2), \underbrace{h_{12}}_{\in \mathcal{W}_r^2}(x_3, x_4) \right)$$



# Shallow Networks

## ① Shallow network: One hidden layer

- Let  $\mathcal{S}_{N,n}$  be the class of shallow networks with  $N$  neurons and  $n$ -dimensional input of the form

$$x \mapsto \sum_{k=1}^N a_k \sigma(w_k \cdot x + b_k), \quad w_k \in \mathbb{R}^n, \quad a_k, b_k \in \mathbb{R},$$

- where  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  is the (nonlinear) activation function.
- Number of trainable parameters:  $p = (n + 2)N \sim N$

## ② Deep neural networks: at least two hidden layers

- Let  $\mathcal{D}_{N,2}$  be the class of Deep networks that take  $n$ -dimensional input with  $N$  neurons and binary tree structure
- Each constituent function in the deep network is a  $\mathcal{S}_{Q,2}$  network
- Complexity:  $(n - 1)Q = N$

# Comparison theorems

## Theorem – Shallow neural networks

Under some assumptions on the non-linear activations, for any  $f \in \mathcal{W}_r^n$  the complexity of shallow networks that provide accuracy at least  $\epsilon$  is  $N = \mathcal{O}(\epsilon^{-n/r})$  and is the best possible  $\implies$  complexity increases **exponentially** with the input dimension - **curse of dimensionality**

## Theorem – Deep neural networks

Under some assumptions on the non-linear activations, for any  $f \in \mathcal{W}_r^{n,2}$  the complexity of a deep network (with the same compositional architecture) that provide accuracy at least  $\epsilon$  is  $N = \mathcal{O}((n-1)\epsilon^{-2/r})$  and is the best possible  $\implies$  **The complexity does NOT** increase exponentially with the input dimension!

# Why is this an interesting result?

- Many functions have an underlying hierarchical composition.  
Example:




$$\begin{aligned}f(x_1, x_2, x_3, x_4) &= ac^2x_4^4x_1^3 + 2acx_1^3x_4^2x_3^3 + ax_1^3x_3^6 \\&+ bc^2x_1x_2x_4^4 + 2bcx_1x_2x_4^4x_3^3 + bx_1x_2x_3^6 \\&= (ax_1^3 + bx_2x_1)(x_3^3 + cx_4^2)^2 \\&= h(h_{11}(x_1, x_2), h_{12}(x_3, x_4))\end{aligned}$$

where



$$h(x, y) = xy^2, \quad h_{11}(x, y) = ax^3 + byx, \quad h_{12}(x, y) = x^3 + cy^2$$

$\implies$  instead of approximating a 9th degree polynomial, a deep network approximates 3 polynomials of 3rd degree

# References I

-  Cybenko, G. (1989). Approximation by superpositions of a sigmoidal function [Publisher: Springer]. *Mathematics of control, signals and systems*, 2(4), 303–314.  
<https://doi.org/10.1007/BF02551274>
-  Hornik, K., Stinchcombe, M., & White, H. (1990). Universal approximation of an unknown mapping and its derivatives using multilayer feedforward networks. *Neural Networks*, 3(5), 551–560.  
[https://doi.org/10.1016/0893-6080\(90\)90005-6](https://doi.org/10.1016/0893-6080(90)90005-6)
-  Mhaskar, H., Liao, Q., & Poggio, T. (2017). When and why are deep networks better than shallow ones? [Number: 1]. *Proceedings of the AAAI Conference on Artificial Intelligence*, 31(1).

# References II

-  Mhaskar, H., & Poggio, T. (2016). Deep vs. shallow networks : An approximation theory perspective. *arXiv:1608.03287 [cs, math]*.
-  Poggio, T., Mhaskar, H., Rosasco, L., Miranda, B., & Liao, Q. (2017). Why and when can deep – but not shallow – networks avoid the curse of dimensionality: A review. *arXiv:1611.00740 [cs]*.